

Introduction to Machine Learning Modeling, Training and Evaluation

Evan Misshula

June 8, 2025

gg

End to end process

Recall **ML workflow** is a sequence of steps to build and deploy a model that solves a problem using data.

The pipeline

:BEAMER_{env}: block :END:

Ingestion & Preprocessing	Analysis	Modeling	Deployment
Definition	EDA	Selection	Tuning
Data Collection	Feature Engineering	Training	Deployment
Cleaning		Evaluation	Monitoring

ML Workflow Graph



Figure: ML workflow steps rendered as a flowchart



What is Model Training?

- Model training is the process of estimating parameters θ of a model $f_{\theta}(x)$ using data $\{(x_i, y_i)\}_{i=1}^n$.
- Typically achieved by minimizing a loss function:

$$\hat{\theta} = \arg \min_{\theta} \frac{1}{n} \sum_{i=1}^n \mathcal{L}(f_{\theta}(x_i), y_i) \quad (1)$$

- Common loss functions:
 - **Squared error loss** (regression): $\mathcal{L}(\hat{y}, y) = (\hat{y} - y)^2$
 - **Cross-entropy loss** (classification):

$$\mathcal{L}(\hat{y}, y) = - \sum_c \mathbb{1}_{\{y=c\}} \log \hat{p}_c \mathbb{1}_{\{x=1\}} \quad (2)$$

Training vs Generalization

- **Empirical risk** (training error):

$$\hat{R}(\theta) = \frac{1}{n} \sum_{i=1}^n \mathcal{L}(f_{\theta}(x_i), y_i) \quad (3)$$

- **Expected risk** (true/generalization error):

$$R(\theta) = \mathbb{E}_{(x,y) \sim \mathcal{D}} [\mathcal{L}(f_{\theta}(x), y)] \quad (4)$$

- Generalization gap: $R(\theta) - \hat{R}(\theta)$
- Overfitting: small \hat{R} , large R

Evaluation Metrics

- Regression:

- Mean Squared Error (MSE):

$$\text{MSE} = \frac{1}{n} \sum_{i=1}^n (\hat{y}_i - y_i)^2$$

- R^2 score:

$$R^2 = 1 - \frac{\sum_i (\hat{y}_i - y_i)^2}{\sum_i (y_i - \bar{y})^2}$$

- Classification:

- Accuracy: $\text{Accuracy} = \frac{1}{n} \sum_{i=1}^n \mathbb{1}_{\{\hat{y}_i = y_i\}}$
- Precision: $\frac{\text{TP}}{\text{TP} + \text{FP}}$
- Recall: $\frac{\text{TP}}{\text{TP} + \text{FN}}$
- F1 score: harmonic mean of precision and recall

$$F1 = 2 \cdot \frac{\text{Precision} \cdot \text{Recall}}{\text{Precision} + \text{Recall}}$$

Cross-Validation

- Cross-validation estimates generalization error by partitioning data.
- **k-fold CV:**
 - Split data into k disjoint subsets.
 - For each $i = 1, \dots, k$:
 - Train on $k - 1$ folds
 - Evaluate on fold i
 - Average the evaluation metrics.

Bias-Variance Tradeoff

- Expected prediction error at point x :

$$\mathbb{E}[(f(x) - y)^2] = \underbrace{[\mathbb{E}(f(x)) - y]^2}_{\text{Bias}^2} + \underbrace{\mathbb{E}[(f(x) - \mathbb{E}(f(x)))^2]}_{\text{Variance}} + \underbrace{\sigma^2}_{\text{Irreducible error}}$$

- Simple models: low variance, high bias
- Complex models: low bias, high variance

Model Selection

- Choose the best model using a **validation set** or **cross-validation**.
- Avoid tuning hyperparameters using the test set.
- Balance:
 - Training error
 - Generalization performance
 - Computational cost

Summary Training and Evaluation

- Training minimizes empirical loss.
- Evaluation uses test or validation data.
- Use metrics appropriate for the task.
- Cross-validation provides robust error estimates.
- The bias-variance tradeoff is fundamental in choosing models.